

Issue 2: Concurrent Programming  
· August 2013

# Concurrent Programming Challenges

By **Florian Kugler**

NEW ON OBJC.IO

## Swift Talk

A weekly video series on Swift programming. [Learn more](#)

### In this article

#### [Concurrency APIs on OS X and iOS](#)

Threads

Grand Central Dispatch

Operation Queues

Run Loops

#### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

[Concurrency](#) describes the concept of running several tasks at the same time. This can either happen in a [time-shared](#) manner on a single CPU core, or truly in parallel if multiple CPU cores are available.

OS X and iOS provide several different APIs to enable concurrent programming. Each of these APIs has different capabilities and limitations, making them suitable for different tasks. They also sit on very different levels of abstraction. We have the possibility to operate very close to the metal, but this also comes with great responsibility to get things right.

Concurrent programming is a very difficult subject with many intricate problems and pitfalls, and it's easy to forget this while using APIs like Grand Central Dispatch or NSOperationQueue. In this article, we'll give you an overview of the different concurrency APIs available on OS X and iOS, and dive deeper into the inherent challenges of concurrent programming, which are independent of the specific APIs you use.

## Concurrency APIs on OS X and iOS

Apple's mobile and desktop operating systems provide a variety of concurrent programming APIs. In this article, we'll cover pthread and NSThread, Grand Central Dispatch, Operation Queues, and NSRunLoop. Technically, run loops are not concurrency APIs, but because they don't enable true parallelism, we'll cover them enough to the topic that it's worth mentioning.

We'll start with the lower-level APIs and work our way up to the higher-level ones. We chose this route because the higher-level APIs are built on top of the lower-level APIs. However, when choosing an API for your use case, you should consider them in the exact opposite order: choose the highest level abstraction that gets the job done and keep your concurrency model very simple.

If you're wondering why we are so persistent recommending high-level abstractions and very simple concurrency code, you should read the second part of this article, [challenges of concurrent programming](#), as well as [Peter Steinberger's thread safety article](#).

### Threads

**Threads** are subunits of processes, which can be scheduled

#### In this article

##### [Concurrency APIs on OS X and iOS](#)

[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)

##### [Challenges of Concurrent Programming](#)

[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

concurrency APIs are built on top of threads under the hood – that’s true for both Grand Central Dispatch and operation queues.

Multiple threads can be executed at the same time (or at least perceived as at the same time) by dividing up small slices of computing time to each thread. If multiple tasks are executed at the same time and are available, then multiple threads can be executed at the same time, therefore lessening the total time.

You can use the [CPU strategy view](#) to see how your code or the framework code is being executed on multiple CPU cores.

The important thing to keep in mind is that when your code gets scheduled for execution, it will be paused in order for another thread to execute. This kind of thread scheduling is a very complex task, and it comes with great complexity, which

### In this article

#### [Concurrency APIs on OS X and iOS](#)

[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)

#### [Challenges of Concurrent Programming](#)

[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

Leaving this complexity aside for a moment, you can either use the [POSIX thread](#) API, or the Objective-C wrapper around this API, `NSThread`, to create your own threads. Here’s a small sample that finds the minimum and maximum in a set of 1 million numbers using `pthread`. It spawns off 4 threads that run in parallel. It should be obvious from this example why you wouldn’t want to use pthreads directly.

OBJECTIVE-C

[SELECT ALL](#)

```
#import <pthread.h>

struct threadInfo {
    uint32_t * inputValues;
    size_t count;
};
```

```

void * findMinAndMax(void *arg)
{
    struct threadInfo const
    uint32_t min = UINT32_MAX;
    uint32_t max = 0;
    for (size_t i = 0; i < info->count; ++i) {
        uint32_t v = info->inputValues[i];
        min = MIN(min, v);
        max = MAX(max, v);
    }
    free(arg);
    struct threadResult * result = malloc(sizeof * result);
    result->min = min;
    result->max = max;
    return result;
}

int main(int argc, const char * argv[]) {
    size_t const count = 100;
    uint32_t inputValues[count];

    // Fill input values with random numbers
    for (size_t i = 0; i < count; ++i) {
        inputValues[i] = arc4random_uniform(UINT32_MAX);
    }

    // Spawn 4 threads to find the minimum and maximum:
    size_t const threadCount = 4;
    pthread_t tid[threadCount];
    for (size_t i = 0; i < threadCount; ++i) {
        struct threadInfo * const info = malloc(sizeof * info);
        size_t offset = (count / threadCount) * i;
        info->inputValues = inputValues + offset;
        info->count = MIN(count - offset, count / threadCount);
        int err = pthread_create(&tid[i], NULL, &findMinAndMax, info);
        NSCAssert(err == 0, @"pthread_create() failed: %d", err);
    }

    // Wait for the threads to exit:
    struct threadResult * results[threadCount];
    for (size_t i = 0; i < threadCount; ++i) {
        int err = pthread_join(tid[i], (void **) &results[i]);
        NSCAssert(err == 0, @"pthread_join() failed: %d", err);
    }

    // Find the min and max:
    void * arg = malloc(sizeof * arg);
    memcpy(arg, results, threadCount * sizeof * results);
    return 0;
}

```

## In this article

### Concurrency APIs on OS X and iOS

Threads

Grand Central Dispatch

Operation Queues

Run Loops

### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

```
max = MAX(max, results[i]->max);  
free(results[i]);  
results[i] = NULL;  
}  
  
NSLog(@"min = %u", min);  
NSLog(@"max = %u", max);  
return 0;  
}
```

NSThread is a simple Objective-C class. In the code look more familiar in a C++ style, we define a thread as a subclass of NSThread. If you want to run in the background, you can define an NSThread subclass like

## In this article

### Concurrency APIs on OS X and iOS

Threads

Grand Central Dispatch

Operation Queues

Run Loops

Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

s  
an  
e

## OBJECTIVE-C

[SELECT ALL](#)

```
@interface FindMinMaxThread
@property (nonatomic) NSUInteger
@property (nonatomic) NSUInteger
- (instancetype)initWithNumbers: (NSArray *)numbers;
@end

@implementation FindMinMaxThread
    NSArray *_numbers;

- (instancetype)initWithNumbers: (NSArray *)numbers
{
    self = [super init];
    if (self) {
        _numbers = numbers;
    }
    return self;
}

- (void)main
{
    NSInteger min;
    NSInteger max;
    // process the data
    self.min = min;
    self.max = max;
}
@end
```

## In this article

[Concurrency APIs on OS X and iOS](#)[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)[Challenges of Concurrent Programming](#)[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

To start new threads, we need to create new thread objects and call their start methods:

## OBJECTIVE-C

SELECT ALL

```

NSMutableDictionary *threads = [NSMutableDictionary new];
NSUInteger numberCount = 0;
NSUInteger threadCount = 4;
for (NSUInteger i = 0; i < threadCount; i++) {
    NSUInteger offset = (NSUInteger)arc4random_uniform((uint32_t)numberCount);
    NSUInteger count = MIN(numberCount - offset, threadCount - i);
    NSRange range = NSMakeRange(i, count);
    NSArray *subset = [self.arrayObjects subarrayWithRange:range];
    FindMinMaxThread *thread = [[FindMinMaxThread alloc] initWithArray:subset];
    [threads addObject:thread];
    [thread start];
}

```

## In this article

## Concurrency APIs on OS X and iOS

Threads

Grand Central Dispatch

Operation Queues

Run Loops

## Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

Now we could observe the threads and see that all our newly spawned threads have finished. We will leave this exercise to the interested reader. One problem is that working directly with threads using the `NSThread` APIs, is a relatively cumbersome mental model of coding very well.

One problem that can arise from directly using threads is that the number of active threads increases exponentially if both your code and underlying framework code spawn their own threads. This is actually a quite common problem in big projects. For example, if you create eight threads to take advantage of eight CPU cores, and the framework code you call into from these threads does the same (as it doesn't know about the threads you already created), you can quickly end up with dozens or even hundreds of threads. Each part of the code involved acted responsibly in itself; nevertheless, the end result is problematic. Threads don't come for free. Each thread ties up memory and kernel resources.

Next up, we'll discuss two queue-based concurrency APIs: Grand Central Dispatch and operation queues. They alleviate this problem by centrally

# Grand Central Dispatch

Grand Central Dispatch (GCD) was introduced in order to make it easier for developers to utilize multiple numbers of CPU cores in consumer devices. We'll talk about GCD in our [article about low-level concurrency](#).

With GCD you don't interact with threads directly. You add blocks of code to queues, and the system decides on which processor to execute them. GCD decides on which processor going to be executed on, and it manages the available system resources. This abstraction hides the threads being created, because the developer is abstracted away from application-level concurrency.

The other important change with GCD is the way about work items in a queue rather than threads. This of concurrency is easier to work with.

GCD exposes five different queues: the main queue running on the main thread, three background queues with different priorities, and one background queue with an even lower priority, which is I/O throttled. Furthermore, you can create custom queues, which can either be serial or concurrent queues. While custom queues are a powerful abstraction, all blocks you schedule on them will ultimately trickle down to one of the system's global queues and its thread pool(s).

## In this article

### [Concurrency APIs on OS X and iOS](#)

Threads

Grand Central Dispatch

Operation Queues

Run Loops

### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

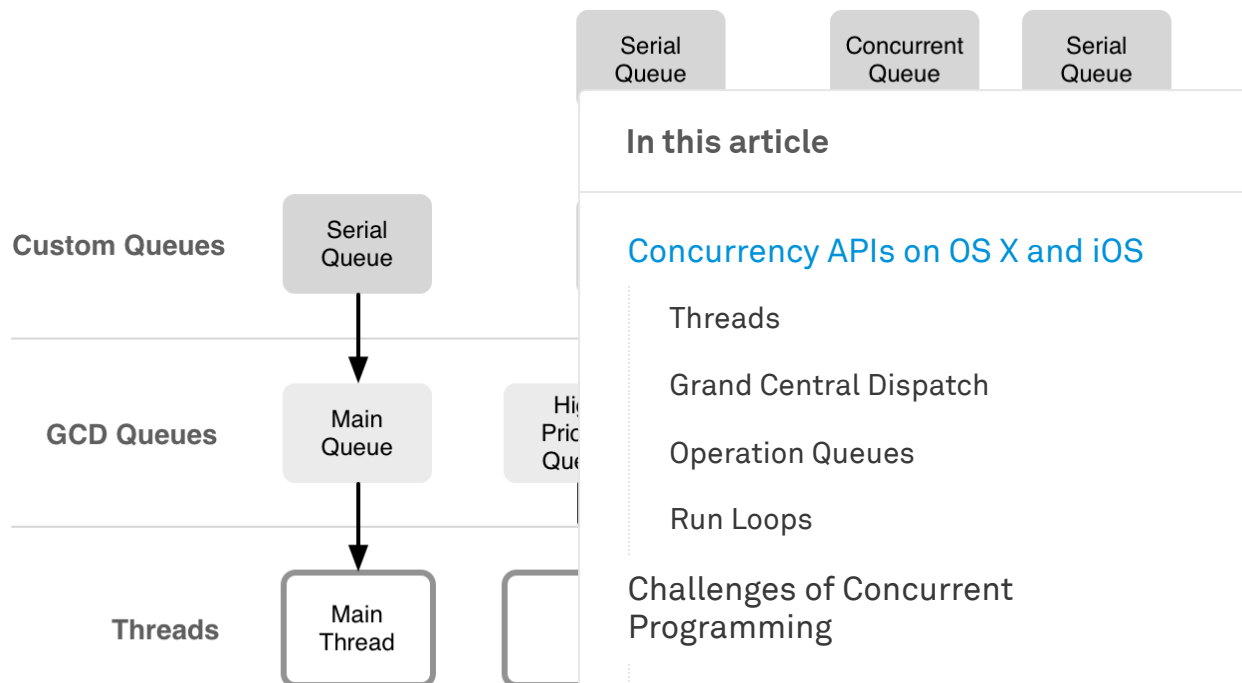
Dead Locks

Starvation

Priority Inversion

Conclusion





Making use of several queues with different priorities can be straightforward at first. However, with a default priority queue in almost all cases, tasks with different priorities can quickly access shared resources. This can cause a program to come to a grinding halt because some low-priority tasks are blocking a high-priority task from executing. You can read more about this phenomenon, called priority inversion, [below](#).

Although GCD is a low-level C API, it's pretty straightforward to use. This makes it easy to forget that all caveats and pitfalls of concurrent programming still apply while dispatching blocks onto GCD queues. Please make sure to read about the [challenges of concurrent programming](#) below, in order to be aware of the potential problems. Furthermore, we have an excellent [walkthrough of the GCD API](#) in this issue that contains many in-depth explanations and valuable hints.

## Operation Queues

implement several convenient features on top of it, which often makes it the best and safest choice for application developers.

The `NSOperationQueue` class has queue and custom queues. The main custom queues are processed in the `main` thread, which are processed by these queues. `NSOperation`.

You can define your own operation `main`, or by overriding `start`. The `start` method gives you less flexibility. In return, the `start` method is finished when `main` returns.

OBJECTIVE-C

SELECT

```
@implementation YourOperation
- (void)main
{
    // do your work here ...
}
@end
```

### In this article

#### [Concurrency APIs on OS X and iOS](#)

Threads

Grand Central Dispatch

Operation Queues

Run Loops

#### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

If you want more control and to maybe execute an asynchronous task within the operation, you can override `start`:

OBJECTIVE-C

SELECT ALL

```

@implementation YourOperation
- (void)start
{
    self.isExecuting = YES;
    self.isFinished = NO;
    // start your work,
}

- (void)finished
{
    self.isExecuting = NO;
    self.isFinished = YES;
}
@end

```

## In this article

## Concurrency APIs on OS X and iOS

- Threads
- Grand Central Dispatch
- Operation Queues
- Run Loops

## Challenges of Concurrent Programming

- Sharing of Resources
- Mutual Exclusion
- Dead Locks
- Starvation
- Priority Inversion

Conclusion

Notice that in this case, you have to manually change the state properties. In order for an operation to change, the state properties have to be changed manually. So make sure to send proper messages to them via default accessor methods.

In order to benefit from the cancellation feature exposed by operation queues, you should regularly check the `isCancelled` property for longer-running operations:

OBJECTIVE-C

SELECT ALL

```

- (void)main
{
    while (notDone && !self.isCancelled) {
        // do your processing
    }
}

```

## OBJECTIVE-C

```

NSOperationQueue *queue = [[
YourOperation *operation = [
[queue addOperation:operati

```

Alternatively, you can also add blocks directly to the queue, which is often handy, e.g. if you want to schedule a block of code to run later.

## OBJECTIVE-C

```

[[NSOperationQueue mainQueue
    // do something...
}];

```

While this is a very convenient way to schedule operations, it can be difficult to debug. If you override the `operationIsExecuting` method, you can easily identify all the operations currently scheduled in a certain queue.

Beyond the basics of scheduling operations or blocks, operation queues offer some features which would be non-trivial to get right in GCD. For example, you can easily control how many operations of a certain queue may be executed concurrently with the `maxConcurrentOperationCount` property. Setting it to one gives you a serial queue, which is great for isolation purposes.

Another convenient feature is the sorting of operations within a queue according to their priorities. This is not the same as GCD's queue priorities. It solely influences the execution order of all operations scheduled in one queue. If you need more control over the sequence of execution beyond the five standard priorities, you can specify dependencies between operations like this:

[SELECT ALL](#)

## In this article

[Concurrency APIs on OS X and iOS](#)[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)[Challenges of Concurrent Programming](#)[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

## OBJECTIVE-C

```
[intermediateOperation addDe  
[intermediateOperation addDe  
[finishedOperation addDepend
```

This simple code guarantees that  
executed before `intermediateOp`  
executed before `finishedOperat`  
powerful mechanism to specify a v  
you create things like operation gr  
executed before the dependent op  
otherwise concurrent queue.

By the very nature of abstractions,  
performance hit compared to using  
cases, this impact is negligible and  
choice.

[SELECT ALL](#)

## In this article

[Concurrency APIs on OS X and iOS](#)[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)[Challenges of Concurrent  
Programming](#)[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

## Run Loops

Run loops are not technically a concurrency mechanism like GCD or operation queues, because they don't enable the parallel execution of tasks. However, run loops tie in directly with the execution of tasks on the main dispatch/operation queue and they provide a mechanism to execute code asynchronously.

Run loops can be a lot easier to use than operation queues or GCD, because you don't have to deal with the complexity of concurrency and still get to do things asynchronously.

A run loop is always bound to one particular thread. The main run loop

kernel events. Whenever you schedule a timer, use a `NSURLConnection` or call `performSelector:withObject:afterDelay:` the run loop is used behind the scenes in order to

Whenever you use a method which remember that run loops can be run defines a set of events the run loop to temporarily prioritize certain tasks

A typical example of this is scrolling loop is not running in its default mode react to, for example, a timer you have stops, the run loop returns to the code been queued up are executed. If you you need to add it to the run loop in

The main thread always has the main threads though don't have a run loop a run loop for other threads too, but the time it is much easier to use the main run loop. If you need to do heavier work that you don't want to execute on the main thread, you can still dispatch it onto another queue after your code is called from the main run loop. Chris has some good examples of this pattern in his article about [common background practices](#).

If you really need to set up a run loop on another thread, don't forget to add at least one input source to it. If a run loop has no input sources configured, every attempt to run it will exit immediately.

### In this article

#### [Concurrency APIs on OS X and iOS](#)

Threads

Grand Central Dispatch

Operation Queues

Run Loops

#### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

## Challenges of Concurrent Programming

Writing concurrent programs comes with many pitfalls. As soon as you're

parallel. Problems can occur in a non-deterministic way, which makes it even more difficult to debug concurrent code.

There is a prominent example for such problems: In 1995, NASA sent the Mars Climate Observer long after a successful landing on Mars, but the mission almost [came to an abrupt end](#). The reasons – it suffered from a phenomenon where a low-priority thread kept blocking a high-priority thread – explore this particular issue in more detail. To demonstrate that even with vast resources available, concurrency can come back to bite you.

### In this article

#### [Concurrency APIs on OS X and iOS](#)

[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)[Challenges of Concurrent Programming](#)[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

## Sharing of Resources


The root of many concurrency related problems is the sharing of resources from multiple threads. A resource can be an object, memory in general, a network connection, etc.

Any resource that is shared between multiple threads is a potential point of conflict, and you have to take safety measures to prevent these kind of conflicts.

In order to demonstrate the problem, let's look at a simple example of a resource in the form of an integer property which you're using as a counter. Let's say we have two threads running in parallel, A and B, and both try to increment the counter at the same time. The problem is that what you write as one statement in C or Objective-C is mostly not just one machine instruction for the CPU. To increment our counter, the current value has to be read from memory. Then the value is incremented by one and finally written back to memory.

Imagine the hazards that can happen if both threads try to do this simultaneously. For example, thread A and thread B both read the value of the counter, increment it, and then write it back to memory. Thread A's increment is lost.

time, thread B also increments the counter by one and writes a 18 back to memory, just after thread A. At this point the data has become corrupted, because the counter has been incremented twice from a 17.



## In this article

### Concurrency APIs on OS X and iOS

- Threads
- Grand Central Dispatch
- Operation Queues
- Run Loops
- Challenges of Concurrent Programming
- Sharing of Resources
- Mutual Exclusion
- Dead Locks
- Starvation
- Priority Inversion

## Conclusion

This problem is called a [race condition](#) where multiple threads access a shared resource.

One thread must be finished operating on a resource before another one begins accessing it. If you're not only writing a simple integer but a more complex structure to memory, it might even happen that a second thread tries to read from this memory while you're in the midst of writing it, therefore seeing half new and half old or uninitialized data. In order to prevent this, multiple threads need to access shared resources in a mutually exclusive way.

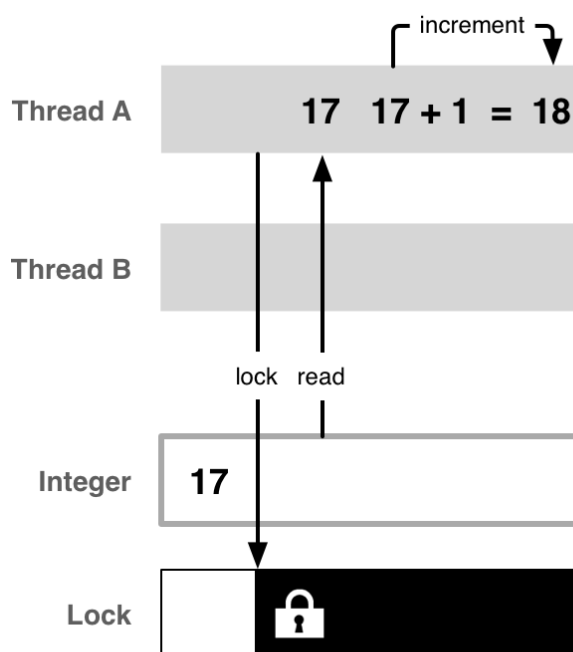
In reality, the situation is even more complicated than this, because modern CPUs change the sequence of reads and writes to memory for optimization purposes ([Out-of-order execution](#)).

## Mutual Exclusion

[Mutual exclusive](#) access means that only one thread at a time gets access



finished its operation, it releases the lock, so that other threads get a chance to access it.



### In this article

#### [Concurrency APIs on OS X and iOS](#)

- Threads
- Grand Central Dispatch
- Operation Queues
- Run Loops

#### Challenges of Concurrent Programming

- Sharing of Resources
- Mutual Exclusion
- Dead Locks
- Starvation
- Priority Inversion

#### Conclusion

In addition to ensuring mutual exclusion, the problem caused by out-of-order CPU accessing the memory in the sequence defined by your program instructions, guaranteeing mutually exclusive access alone is not enough. To work around this side effect of CPU optimization strategies, [memory barriers](#) are used. Setting a memory barrier makes sure that no out-of-order execution takes place across the barrier.

Of course the implementation of a mutex lock in itself needs to be race-condition free. This is a non-trivial undertaking and requires use of special instructions on modern CPUs. You can read more about atomic operations in Daniel's [low-level concurrency techniques](#) article.

Objective-C properties come with language level support for locking in the form of declaring them as atomic. In fact, properties are even atomic by default. Declaring a property as atomic results in implicit locking/unlocking around each access of this property. It might be

Acquiring a lock on a resource always comes with a performance cost. Acquiring and releasing a lock needs to be race-condition free, which is non-trivial on multi-core systems. A thread that acquires a lock might have to wait because some other thread already holds it. In this case, that thread will sleep and the thread that acquired the lock relinquishes the lock. All of this is complicated.

There are different kinds of locks. Some locks have no lock contention but perform poorly. Some locks are more expensive at a base level, but they avoid contention. Contention is the situation when one thread tries to acquire a lock that has already been taken).

There is a trade-off to be made here. Acquiring a lock at a price (lock overhead). Therefore, threads are constantly entering and exiting critical sections (acquiring and releasing locks). At the same time, if you have a large region of code, you run the risk of a deadlock.

Threads are often unable to do work because they're waiting to acquire a lock. It's not an easy task to solve.

It is quite common to see code which is supposed to run concurrently, but which actually results in only one thread being active at a time, because of the way locks for shared resources are set up. It's often non-trivial to predict how your code will get scheduled on multiple cores. You can use Instrument's [CPU strategy view](#) to get a better idea of whether you're efficiently using the available CPU cores or not.

## Dead Locks

Mutex locks solve the problem of race conditions, but unfortunately they

### In this article

#### Concurrency APIs on OS X and iOS

Threads

Grand Central Dispatch

Operation Queues

Run Loops

Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion


Dead Locks

Starvation

Priority Inversion

Conclusion

**locks.** A dead lock occurs when multiple threads are waiting on each other to finish and get stuck.



### In this article

#### Concurrency APIs on OS X and iOS

- Threads
- Grand Central Dispatch
- Operation Queues
- Run Loops

#### Challenges of Concurrent Programming

- Sharing of Resources
- Mutual Exclusion
- Dead Locks
- Starvation
- Priority Inversion

#### Conclusion

Consider the following example code variables:

OBJECTIVE-C [SELECT ALL](#)

```
void swap(A, B)
{
    lock(lockA);
    lock(lockB);
    int a = A;
    int b = B;
    A = b;
    B = a;
    unlock(lockB);
    unlock(lockA);
}
```

This works quite well most of the time. But when by chance two threads call it at the same time with opposite variables

OBJECTIVE-C

[SELECT ALL](#)

```
swap(X, Y); // thread 1
swap(Y, X); // thread 2
```

we can end up in a dead lock. Thread 1 acquires a lock on Y. Now they're both locked and never be able to acquire it.

Again, the more resources you share, the greater your risk of running into one more reason to keep things as simple as possible between threads. See the [section about doing things asynchronously](#) in the [APIs article](#).

## Starvation

Just when you thought that there are enough problems to think of, a new one comes around the corner. Locking shared resources can result in the [readers-writers problem](#). In many cases, it would be wasteful to restrict reading access to a resource to one access at a time. Therefore, taking a reading lock is allowed as long as there is no writing lock on the resource. In this situation, a thread that is waiting to acquire a write lock can be starved by more read locks occurring in the meantime.

In order to solve this issue, more clever solutions than a simple read/write lock are necessary, e.g. giving [writers preference](#) or using the [read-copy-update](#) algorithm. Daniel shows in his [low-level concurrency techniques](#) article how to implement a multiple reader/single writer pattern with GCD which doesn't suffer from writer starvation.

### In this article

#### Concurrency APIs on OS X and iOS

[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)

#### Challenges of Concurrent Programming

[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)

# Priority Inversion

We started this section with the example of Mars suffering from a concurrency problem. We look why Pathfinder almost failed, and how we can avoid it from the same problem, called **priority inversion**.

Priority inversion describes a condition where a higher priority task from executing because a lower priority task holds a resource. Since GCD exposes background queues, one which even is I/O throttled, it's easy to see how this can happen.

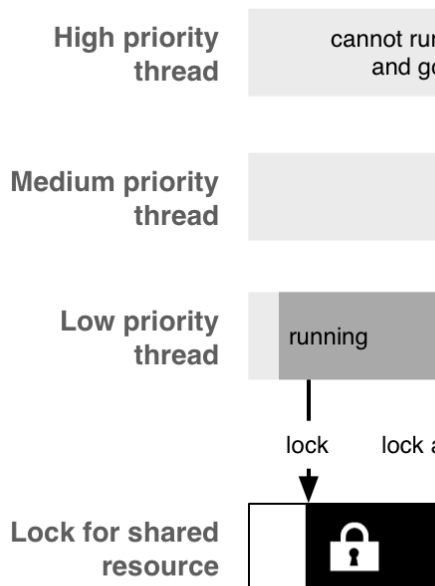
The problem can occur when you have two tasks that share a common resource. When the low-priority task holds the common resource, it is supposed to release its lock and to let the high-priority task to execute, causing significant delays. Since the high-priority task is waiting, it's called priority inversion.

As long as the low-priority task has the lock, there is a window of opportunity for medium-priority tasks to run and to preempt the low-priority task, because the medium-priority tasks have now the highest priority of all currently runnable tasks. At this moment, the medium-priority tasks hinder the low-priority task from releasing its lock, therefore effectively gaining priority over the still waiting, high-priority tasks.

## In this article

### [Concurrency APIs on OS X and iOS](#)

[Threads](#)[Grand Central Dispatch](#)[Operation Queues](#)[Run Loops](#)[Challenges of Concurrent Programming](#)[Sharing of Resources](#)[Mutual Exclusion](#)[Dead Locks](#)[Starvation](#)[Priority Inversion](#)[Conclusion](#)



## In this article

### Concurrency APIs on OS X and iOS

- Threads
- Grand Central Dispatch
- Operation Queues
- Run Loops

### Challenges of Concurrent Programming

- Sharing of Resources
- Mutual Exclusion
- Dead Locks
- Starvation
- Priority Inversion

### Conclusion

In your own code, things might not have occurred in the Mars rover, as priority inversion is a less severe manner.

In general, don't use different priorities. If you're using different priorities, more likely than not, it's actually going to make things worse. always use the default priority queue (directly, or as a target queue). If you're using different priorities, more likely than not, it's actually going to make things worse.

The lesson from this is that using multiple queues with different priorities sounds good on paper, but it adds even more complexity and unpredictability to concurrent programs. And if you ever run into a weird problem where your high-priority tasks seem to get stuck for no reason, maybe you will remember this article and the problem called priority inversion, which even the NASA engineers encountered.

## Conclusion

resulting behavior quickly gets very difficult to oversee, and debugging these kind of problems is often very hard.

On the other hand, concurrency is a natural extension of the computing power of modern multi-processor systems. A concurrency model as simple as possible can reduce the amount of locking necessary.

A safe pattern we recommend is to do all the work on the main thread, then use an operation queue to move the work to the background, and finally get back the result of your background work. This pattern is safe for yourself, which greatly reduces the

### In this article

#### Concurrency APIs on OS X and iOS

Threads

Grand Central Dispatch

Operation Queues

Run Loops

#### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion

**CONTINUE READING ISSUE 2**

## Concurrent Programming

July 2013

### Editorial

#### → Concurrent Programming: APIs and Challenges

by Florian Kugler

#### Common Background Practices

by Chris Eidhof

#### Low-Level Concurrency APIs

by Daniel Eggert

#### Thread-Safe Class Design

by Peter Steinberger



SWIFT TALK

BOOKS

WORKSHOPS



by Tobias Krüntzer

EXPLORE THE ARCHIVE

## Year 1

- #1 [Lighter View Controllers](#)
- #2 [Concurrent Programming](#)
- #3 [Views](#)
- #4 [Core Data](#)
- #5 [iOS 7](#)
- #6 [Build Tools](#)
- #7 [Foundation](#)
- #8 [Quadcopter Project](#)
- #9 [Strings](#)
- #10 [Syncing Data](#)
- #11 [Android](#)
- #12 [Animations](#)

### In this article

#### Concurrency APIs on OS X and iOS

- Threads
- Grand Central Dispatch
- Operation Queues
- Run Loops

#### Challenges of Concurrent Programming

- Sharing of Resources
- Mutual Exclusion
- Dead Locks
- Starvation
- Priority Inversion

#### Conclusion

#24 [Audio](#)

### OUR LATEST BOOK

## Advanced Swift

A deep dive into Swift's features, from low-level programming to high-level abstractions.

objc ↑↓  
Advanced  
Swift

Updated for Swift 3

BUY NOW

STAY UP TO DATE

objc ↑↓

SWIFT TALK

BOOKS

WORKSHOPS





very occasional updates about upcoming books, promotions, and other things.

Your Email

objc 

objc.io publishes books, videos, and more about iOS and OS X development.

LEARN

Swift Talk

Books

Workshops

Issues

FOLLOW

Blog

Newsletters

Twitter

YouTube

## In this article

### Concurrency APIs on OS X and iOS

Threads

Grand Central Dispatch

Operation Queues

Run Loops

### Challenges of Concurrent Programming

Sharing of Resources

Mutual Exclusion

Dead Locks

Starvation

Priority Inversion

Conclusion